

## CSCI 416 - HW3

Name: Nicholas Wilson

### Problem 3

```
"""
```

```
=====
Test SVM with custom Gaussian kernels
=====
```

*Author: Eric Eaton, 2014*

*Adapted from scikit\_learn documentation.*

```
"""
```

```
print(__doc__)

from numpy import loadtxt, ones, zeros, where
import numpy as np
import matplotlib.pyplot as plt
from sklearn import svm, datasets
from svmKernels import myGaussianKernel
from svmKernels import _gaussSigma

# import some data to play with

filename = 'data/svmTuningData.dat'
data = loadtxt(filename, delimiter=',')
X_raw = data[:, 0:-1]
Y_raw = np.squeeze(np.array([data[:, 2]]).T)
m, d = X_raw.shape
#print(m,d)

print("Training the SVMs...")

trials = 10
folds = 10
fold_size = m//folds
train_size = fold_size * (folds - 1)
#print(fold_size, train_size)
validation_size = fold_size
```

```

# Best parameters
best=[0,0,0]

# Search parameters through a grid
#TO DO: CHANGE THIS PART

sigma_vals = 10**np.arange(-3., 5.)
C_vals = np.linspace(1,100, num=100)

for c in range(0,100):
    for g in range(0,8):
        C = C_vals[c]
        _gaussSigma = sigma_vals[g]
        accuracy = 0

        for t in range(0,trials):

            #randomize the data set
            p = np.random.permutation(m)
            order = p[0:m]
            X = X_raw[order,:]
            Y = Y_raw[order]

            for f in range(0,folds):

                #cross validation: get train set and test set
                # TO DO
                # You need to finish the following for cross validation
                #Use fold_size and validation_size

                #X_train is Equal to all but the "fth" fold of X
                X_train = np.ones((train_size, d)) #Create an array of the
correct size
                #Grab all the rows of X up to the f * fold_size row
                X_train[0:f * fold_size,:] = X[0:f * fold_size,:]
                #Grab all the rows of X after f * fold_size +
validation_size row
                X_train[f * fold_size:m,:] = X[f * fold_size +
validation_size:m,:]
                # print(X.shape)
                # print(X)
                # print()
                # print(X_train.shape)
                # print(X_train)
                Y_train = np.ones((train_size)) #Equal to all but the
"fth" fold of Y
                Y_train[0:f*fold_size] = Y[0:f*fold_size]
                Y_train[f*fold_size:] = Y[f*fold_size + validation_size:]
                # print(Y.shape)

```

```

        # print(Y)
        # print()
        # print(Y_train.shape)
        # print(Y_train)
        #X_test = np.ones(validation_size, d) #Equal to the "fth"
fold of X
        X_test = X[f * fold_size: f * fold_size +
validation_size,:]
        # print(X_test)
        # print(X_test.shape)
        Y_test =Y[f * fold_size: f * fold_size + validation_size]
#Equal to the "fth" fold of Y
        # print()
        # print(Y_test.shape)
        # print(Y_test)
        # create an instance of SVM with build in RBF kernel and
train it
        equivalentGamma = 1.0 / (2 * _gaussSigma ** 2)
        model = svm.SVC(C=C, kernel='rbf', gamma=equivalentGamma)
        model.fit(X_train, Y_train)

        predictions_test =
model.predict(np.c_[X_test[:,0],X_test[:,1]])
        a = np.mean(Y_test==predictions_test)
        accuracy += a

        # Best Accuracy So Far
        average_accuracy = accuracy/(folds*trials)
        if average_accuracy > best[2]:
            best[0] = C
            best[1] = _gaussSigma
            best[2] = average_accuracy

print(best)

```

```

=====
Test SVM with custom Gaussian kernels
=====

```

Author: Eric Eaton, 2014

Adapted from scikit\_learn documentation.

Training the SVMs...  
[82.0, 10.0, 0.9830769230769227]

```

# create an instance of SVM with build in RBF kernel and train it
print("The best parameters are: ", best)
C = best[0]
_gaussSigma = best[1]
equivalentGamma = 1.0 / (2 * _gaussSigma ** 2)
model = svm.SVC(C=C, kernel='rbf', gamma=equivalentGamma)
model.fit(X, Y)

h = .02 # step size in the mesh

# Plot the decision boundary. For that, we will assign a color to each
# point in the mesh [x_min, m_max]x[y_min, y_max].
x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min,
y_max, h))

predictions = model.predict(np.c_[xx.ravel(), yy.ravel()])
predictions = predictions.reshape(xx.shape)

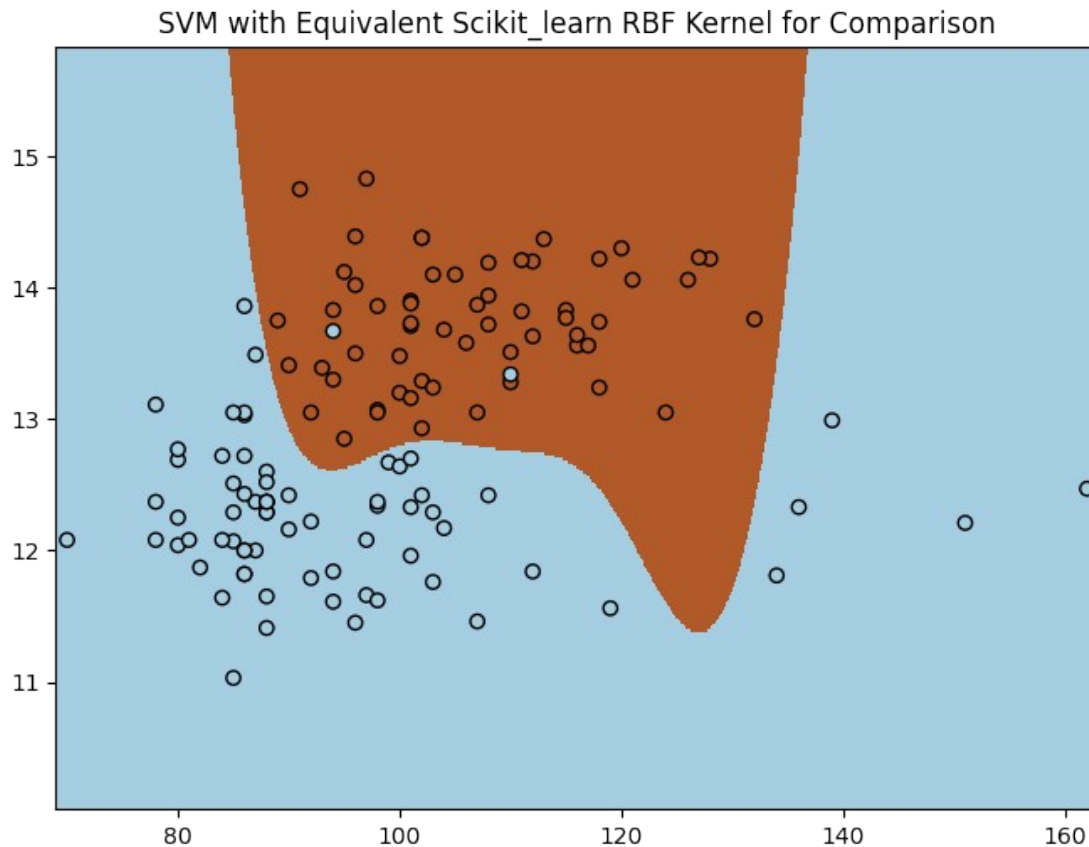
# plot my results
plt.figure(figsize=(8, 6), dpi=100)

plt.pcolormesh(xx, yy, predictions, cmap="Paired")
plt.scatter(X[:, 0], X[:, 1], c=Y, cmap="Paired", edgecolors="black")
# Plot the training points
plt.title('SVM with Equivalent Scikit_learn RBF Kernel for
Comparison')
plt.axis('tight')

plt.show()

The best parameters are: [82.0, 10.0, 0.9830769230769227]

```



Report optimal values and the corresponding estimated accuracy. And explain how you find those optimal values.

My best  $C$  was 82.0 and my best sigma was 10.0. Corresponding accuracy was 0.9830769230769227, or 98.308% accurate. I found those by testing every sigma value between 0.001 and 10000 and every  $C$  value between 1 and 100. In each test, I did 10 trials. Within each trial, I randomized the data, split it into 10 folds and used 9 of those folds for training data and the remaining fold for testing data, and used that to calculate accuracy for that trial. For each  $C$  and sigma value, I then calculated the average accuracy across those 10 trials, then reported the highest average accuracy.

## Problem 4

### Movie Recommendations

user

Alice

Bob

Carol

David

Eve

What movie should I recommend to Bob? Will Carol like Frozen?

**Goal:** Fill in entries of the "rating matrix"

## Problem Setup

Let's formalize this as a machine learning problem. To make it concrete, let's load some data and see what it looks like.

```
%matplotlib inline
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from IPython.display import display
import scipy.io

# Load train and test data
data = scipy.io.loadmat('movies.mat')

titles = [t[0] for t in data['movieData']['title'][0,0].ravel()]

for x,y in data.items():
    if isinstance(y, (np.ndarray)) and len(y)==1:
        data[x] = np.asscalar(y)
    elif isinstance(y, (np.ndarray)):
        data[x] = y.ravel()

nUsers    = data['nUsers']
nMovies   = data['nMovies']
userData  = data['userData']
movieData = data['movieData']

train_user  = data['train_user']-1 # matlab 1-index correction
train_movie = data['train_movie']-1 # matlab 1-index correction
train_rating = data['train_rating']

valid_user  = data['valid_user']-1 # matlab 1-index correction
valid_movie = data['valid_movie']-1 # matlab 1-index correction
valid_rating = data['valid_rating']

test_user   = data['test_user']-1 # matlab 1-index correction
test_movie  = data['test_movie']-1 # matlab 1-index correction
```

```

# Create a pandas data frame for training data to facilitate
# visualization and inspection

train_title = [titles[i] for i in train_movie]

train_data = pd.DataFrame(data = {'user_id' : train_user,
                                  'movie_id' : train_movie,
                                  'rating' : train_rating,
                                  'title': train_title},
                           columns = ['user_id', 'movie_id', 'rating',
                                     'title'])

```

```

# subsample to 5000 rows to more easily see a small sampling of
# ratings for each user
train_data = train_data[:5000]

```

```

# sort by user
train_data = train_data.sort_values(by=['user_id', 'rating'])

```

```
display(train_data)
```

```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:15:
DeprecationWarning: np.asscalar(a) is deprecated since NumPy v1.16,
use a.item() instead
  from ipykernel import kernelapp as app

```

	user_id	movie_id	rating	
title				
2070	0	242	1	Jungle2Jungle
(1997)				
2175	0	73	1	Faster Pussycat! Kill! Kill!
(1965)				
984	0	101	2	Aristocats, The
(1970)				
2400	0	236	2	Jerry Maguire
(1996)				
4364	0	179	3	Apocalypse Now
(1979)				
...	...	...	...	
...				
1373	942	61	3	Stargate
(1994)				
724	942	150	4	Willy Wonka and the Chocolate Factory
(1971)				
1883	942	23	4	Rumble in the Bronx
(1995)				
3403	942	731	4	Dave
(1993)				

1851	942	11	5	Usual Suspects, The
(1995)				

[5000 rows x 4 columns]

## Training Data

As we can see, the training data presents observed entries of the "ratings" matrix as list of triples  $(i_k, j_k, r_k)$  where

- $i_k$  is the user index of  $k$ th rating
- $j_k$  is the movie index of  $k$ th rating
- $r_k$  is the value of  $k$ th rating (1-5)

In our code we will store the entries of the tuples in three separate 1d arrays of the same length, so the  $k$ th rating is represented by the values `train_user[k]`, `train_movie[k]`, and `train_rating[k]`.

## Problem Formulation

Now, let's formulate the problem mathematically. Suppose there are  $m$  users and  $n$  movies. Let  $R$  be the  $m \times n$  "rating" matrix, where  $R_{ij}$  is the (possibly unknown) rating for user  $i$  on movie  $j$ .

Our training data gives us some of the entries of the rating matrix. Our goal is to learn a parametric model to predict entries that we don't observe.

### But Where are the Features?

What sort of predictive model can we use for entries of  $R$ ?

In past learning problems we had *feature vectors* and we learned *weight vectors* to make predictions (using dot products).

Now we do not have feature vectors. What should we do?

## Matrix Factorization Model

Our solution is to **learn weight vectors for both users and movies**.

Let  $\mathbf{u}_i \in \mathbb{R}^d$  be the weight vector for user  $i$  and  $\mathbf{v}_j \in \mathbb{R}^d$  be the weight vector for movie  $j$ . Then we can predict the rating for user  $i$  on movie  $j$  as:

$$H_{ij} = \mathbf{u}_i^T \mathbf{v}_j$$

Our goal is to learn weight vectors for every user and movie so that  $R_{ij} \approx H_{ij}$  for those entries of the rating matrix that we observe.



**Problem statement:** Given observed entries of the rating matrix presented as triples  $(i_k, j_k, r_k)$  for  $k=1, \dots, n_{\text{train}}$ , find weight vectors  $u_i$  for each user  $i$  and  $v_j$  for each movie  $j$  such that:

$$r_k \approx u_{i_k}^T v_{j_k}, k=1, 2, \dots, n_{\text{train}}$$

## Why is This Called Matrix Factorization?

- Place the user weight vectors  $u_i$  into the rows of a matrix  $U$  and the movie feature vectors  $v_j$  into the rows of a matrix  $V$

$$U = \begin{bmatrix} u_1^T \\ u_2^T \\ \vdots \\ u_m^T \end{bmatrix} \in \mathbb{R}^{m \times d} \quad V = \begin{bmatrix} v_1^T \\ v_2^T \\ \vdots \\ v_n^T \end{bmatrix} \in \mathbb{R}^{n \times d}$$

- Consider the product  $UV^T$ :

$$U$$

- It is easy to check that  $(i, j)$  entry of  $UV^T$  is equal to  $u_i^T v_j$ , which is our prediction for the  $(i, j)$  entry of  $R$
- In other words, our model is that  $R \approx UV^T$  (a **factorization** of  $R$ )
- We choose  $U$  and  $V$  to get good predictions for those entries of  $R$  that we can observe. As long as we don't overfit, this gives us power to generalize to entries we don't observe
- The "hidden dimension"  $d$  (the length of each weight vector) is a hyperparameter that must be tuned with hold-out data.

## Your Job: Solve the Learning Problem

- Formulate a squared error cost function corresponding to the problem statement above.
- Add regularization for *every* user weight vector  $u_i$  and movie weight vector  $v_j$  to get a regularized cost function
- Write down the partial derivatives of your regularized cost function with respect to the entries of  $u_i$  and  $v_j$
- Plug the partial derivatives into stochastic gradient descent (SGD) and write down the update rule
- Implement SGD
- Tune parameters (e.g., dimension  $d$ , regularization parameter) get good performance on the validation set

## Logistics

- Submit predictions on test set

- Evaluation: root-mean squared error (RMSE) on test set

$$\text{RMSE} = \sqrt{\frac{1}{n_{\text{test}}} \sum_{(i,j) \in \text{test set}} (H_{ij} - R_{ij})^2}$$

- Your grade:

RMSE

---

$\leq 1.0$

$\leq 0.97$

$\leq 0.95$

$\leq 0.94$

### (Review on your own) Model Extension: Add Biases

To get really great performance, consider this extended model for a predicted rating:

$$H_{ij} = \mu + a_i + b_j + u_i^T v_j$$

This adds several terms to the prediction for user  $i$  on movie  $j$ :

- $\mu$  is an overall baseline rating. For example, the overall average rating of all users on all movies may be  $\mu = 3.3$
- $a_i$  is a user-specific adjustment or "bias". For example, perhaps Alice really loves movies and gives them all high ratings. Then, her bias might be  $a_i = +0.4$ . But Bob is hard to please, so his bias is  $a_i = -0.7$ .
- $b_j$  is a movie-specific bias. For example, perhaps Inside Out is universally loved, so its bias is  $b_j = +0.7$ . A really bad movie would have a negative bias.

The set of parameters of this model includes:

- $\mu$
- $a_i, i = 1, \dots, m$
- $b_j, j = 1, \dots, n$
- $u_i \in \mathbb{R}^d, i = 1, \dots, m$
- $v_j \in \mathbb{R}^d, j = 1, \dots, n$

To learn these parameters, derive partial derivatives of the regularized cost function with respect to *all* of the above parameters, and update them all within your stochastic gradient descent loop.

### Further Reading

[Matrix Factorization Techniques for Recommender Systems](#) by Yehuda Koren, Robert Bell and Chris Volinsky

- Authors were on the winning team of Netflix prize
- Paper includes algorithms---but beware different notation

## Step 0: Familiarize Yourself With Variables

Here are the variables we populated while loading the data above --- make sure you run that cell first.

```
# 1) Metadata
#
#     nUsers      # of users
#     nMovies     # of movies
#     titles      list of movie titles
#
#
# 2) Training data (60K ratings). This consists of three 1d arrays,
#     each of length 60K:
#
#     train_user, train_movie, train_rating
#
#     The entries specify the ratings:
#
#     train_user[k]    user index of kth rating
#     train_movie[k]   movie index of kth rating
#     train_rating[k]  value (1-5) of kth rating
#
# 2) Validation data (20K ratings). Three vectors of length 20K:
#
#     valid_user, valid_movie, valid_rating
#
#     Use this to evaluate your model and tune parameters.
#
# 3) Test set (20K user-movie pairs without ratings):
#
#     test_user, test_movie
#
#     You will create predictions for these pairs and submit them for
#     grading.
```

## Step 1: Look at the Prediction Method

To make things concrete, first take a look at the prediction method below. This is just a stub for now that returns the same value  $\mu$  for every prediction. Later you will update this to make predictions given the weight vectors and biases.

```
from os import X_OK
def rmse(h, r):
    resid = h - r
    cost = np.sqrt(np.mean(resid**2))
    return cost
```

```

def predict(mu, uBiases, mBiases, uWeights, mWeights, user, movie):
    '''
    PREDICT Make predictions for user/movie pairs
    Inputs:
        model parameters
        mu                average user rating for all movies
        uBiases           vector of biases for each user
        mBiases           vector of biases for each movie
        uWeights          matrix of user weights (eg U)
        mWeights          matrix of movie weights (eg V)
        user              vector of users (eg train_user)
        movie             vector of movies (eg train_movie)

    Output:
        predictions       vector of predictions
    '''

    # This is a stub that predicts the mean rating for all user-movie
    pairs
    # Replace with your code.

    L = len(user)
    predictions = np.zeros(L)
    for k in range(L):
        i = user[k]
        x = movie[k]
        predictions[k] = mu + uBiases[i] + mBiases[x] +
np.dot(uWeights[i].T, mWeights[x])
        #predictions[:] = mu

    return predictions

```

## Step 2: Learning and Validation

Write code here to do the learning and validation. Stubs are provided. Make sure you derive the partial derivatives on paper before you try to code them.

```

#####
# Tunable parameters (you will add more)
#####

nDims = 30 #100

#Regularization constants
lambda_U = 5e-3 #Regularization constant for user vectors, 4
lambda_V = 5e-3 #Regularization constant for movie vectors, 15
lambda_a = 1e-4 #Regularization constant for user biases, 0.05
lambda_b = 1e-4 #Regularization constant for movie biases, 0.05
#lambdas = [lambda_U, lambda_V, lambda_a, lambda_b]

```

```

#Alpha/learning rates
alpha_U = 2.5e-3 #User vectors, 1e-2
alpha_V = 2.5e-3 #Movie vectors, 1e-2
alpha_a = 2.5e-3 #User biases, 1e-2
alpha_b = 2.5e-3 #Movie biases, 1e-2
#alphas = [alpha_U,alpha_V,alpha_a,alpha_b]

#Number of iterations
iterations = 1800000

#####
# Initialize parameters
#####

mu = np.mean(train_rating)
a = np.zeros(nUsers)
b = np.zeros(nMovies)
U = np.random.randn(nUsers, nDims) *.01 # User weights
V = np.random.randn(nMovies, nDims) *.01 # Movie features

#####
# Training and validation
#####

# TODO: write code to train model and evaluate performance on
validation set
#
# predict() is a stub that predicts the overall mean for all user-
movie
# pairs. Update it to take more parameters and make real predictions.

#Write the SGD loop HERE
#Use it to modify a, b, U, and V
#THEN use it to make predictions
num_ratings = len(train_rating)
for iter in range(iterations): #Run SGD loop iterations times
    #Select a data point randomly
    k = np.random.randint(0,num_ratings)
    #Get the user and movie corresponding to this data point
    i = train_user[k]
    x = train_movie[k]
    #Calculate epsilon_k
    e_k = 2 * (train_rating[k] - (np.dot(U[i].T, V[x]) + mu + a[i] +
b[x]))
    #Update U[i], V[x], a[i] and b[x]
    U[i] = U[i] + alpha_U * (e_k * V[x] - lambda_U * U[i])
    V[x] = V[x] + alpha_V * (e_k * U[i] - lambda_V * V[x])
    a[i] = a[i] + alpha_a * (e_k - lambda_a * a[i])

```

```
b[x] = b[x] + alpha_b * (e_k - lambda_b * b[x])
```

```
#Now I have a trained set of user vectors U and movie vectors V  
train_predictions = predict(mu, a, b, U, V, train_user, train_movie)  
valid_predictions = predict(mu, a, b, U, V, valid_user, valid_movie)
```

```
train_rmse = rmse(train_predictions, train_rating)  
valid_rmse = rmse(valid_predictions, valid_rating)
```

```
print('train_rmse=%.3f, valid_rmse=%.3f' % (train_rmse, valid_rmse))
```

```
#####  
# Testing  
#####
```

```
# Make and save predictions for test set  
test_predictions = predict(mu, a, b, U, V, test_user, test_movie)  
np.savetxt('test_predictions.txt', test_predictions)
```

```
train_rmse=0.897, valid_rmse=0.938
```

```
#####  
# Tunable parameters (you will add more)  
#####
```

```
nDims = 30 #100
```

```
#Regularization constants  
lambda_U = 5e-3 #Regularization constant for user vectors, 4  
lambda_V = 5e-3 #Regularization constant for movie vectors, 15  
lambda_a = 1e-4 #Regularization constant for user biases, 0.05  
lambda_b = 1e-4 #Regularization constant for movie biases, 0.05  
#lambdas = [lambda_U, lambda_V, lambda_a, lambda_b]
```

```
#Alpha/learning rates  
alpha_U = 2.5e-3 #User vectors, 1e-2  
alpha_V = 2.5e-3 #Movie vectors, 1e-2  
alpha_a = 2.5e-3 #User biases, 1e-2  
alpha_b = 2.5e-3 #Movie biases, 1e-2
```

```
#Number of iterations  
iterations = 50
```

```
#####  
# Initialize parameters  
#####
```

```
mu = np.mean(train_rating)
```

```

a = np.zeros(nUsers)
b = np.zeros(nMovies)
U = np.random.randn(nUsers, nDims) *.01 # User weights
V = np.random.randn(nMovies, nDims) *.01 # Movie features

#####
# Training and validation
#####

# TODO: write code to train model and evaluate performance on
validation set
#
# predict() is a stub that predicts the overall mean for all user-
movie
# pairs. Update it to take more parameters and make real predictions.

#Write the SGD loop HERE
#Use it to modify a, b, U, and V
#THEN use it to make predictions
num_ratings = len(train_rating)
for iter in range(iterations): #Run SGD loop iterations times
    for k in range(num_ratings):
        i = train_user[k]
        x = train_movie[k]
        #Calculate epsilon_k
        e_k = 2 * (train_rating[k] - (np.dot(U[i].T, V[x]) + mu + a[i] +
b[x]))
        #Update U[i], V[x], a[i] and b[x]
        U[i] = U[i] + alpha_U * (e_k * V[x] - lambda_U * U[i])
        V[x] = V[x] + alpha_V * (e_k * U[i] - lambda_V * V[x])
        a[i] = a[i] + alpha_a * (e_k - lambda_a * a[i])
        b[x] = b[x] + alpha_b * (e_k - lambda_b * b[x])

#####
# Testing
#####
train_predictions = predict(mu, a, b, U, V, train_user, train_movie)
valid_predictions = predict(mu, a, b, U, V, valid_user, valid_movie)

train_rmse = rmse(train_predictions, train_rating)
valid_rmse = rmse(valid_predictions, valid_rating)

print('train_rmse=%.5f, valid_rmse=%.5f' % (train_rmse, valid_rmse))

# Make and save predictions for test set
test_predictions = predict(mu, a, b, U, V, test_user, test_movie)
np.savetxt('test_predictions.txt', test_predictions)

```

```
train_rmse=0.74777, valid_rmse=0.92920
```

## Bonus Material: Inspect Predictions for Different Users

After you have learned a good model, you may wish to interpret what it has learned. We can do this by looking at the most positive and most negative predictions for different users (or the movies that are bumped up or down from the baseline the most).

Read and run the code below to see if you can understand the predictions. (Note: the predictions won't make sense until you have learned a good model!)

```
all_movies = range(nMovies)
```

```
def get_lowest(vals):  
    most_negative = np.argsort(vals)  
    return most_negative
```

```
def get_highest(vals):  
    most_negative = np.argsort(vals)  
    most_positive = most_negative[::-1]  
    return most_positive
```

```
k = 8  
all_users = range(nUsers)  
users_to_examine = all_users[0:5]
```

```
for user in users_to_examine:
```

```
    # Changes from baseline movie predictions for this user  
    delta = np.dot(V, U[user,:])
```

```
    print('*** User %d ***' % (user))  
    print('  Top movies')  
    for i in get_highest(delta)[0:k]:  
        print('    %+.4f %s' % (delta[i], titles[i]))  
    print('')
```

```
    print('  Bottom movies')  
    for i in get_lowest(delta)[0:k]:  
        print('    %+.4f %s' % (delta[i], titles[i]))  
    print('')
```

```
*** User 0 ***
```

```
Top movies  
+0.0000 Crude Oasis, The (1995)  
+0.0000 Last Klezmer: Leopold Kozłowski, His Life and Music, The  
(1995)  
+0.0000 Wend Kuuni (God's Gift) (1982)  
+0.0000 Glass Shield, The (1994)  
+0.0000 Scream of Stone (Schrei aus Stein) (1991)
```



+0.0000 8 Heads in a Duffel Bag (1997)  
+0.0000 Lashou shentan (1992)  
+0.0000 Coldblooded (1995)

Bottom movies

-0.0000 Tainted (1998)  
-0.0000 Girls Town (1996)  
-0.0000 Walk in the Sun, A (1945)  
-0.0000 Wings of Courage (1995)  
-0.0000 Other Voices, Other Rooms (1997)  
-0.0000 Silence of the Palace, The (Saint el Qusur) (1994)  
-0.0000 My Favorite Season (1993)  
-0.0000 Substance of Fire, The (1996)

\*\*\* User 1 \*\*\*

Top movies

+0.0000 Other Voices, Other Rooms (1997)  
+0.0000 Leopard Son, The (1996)  
+0.0000 Stranger, The (1994)  
+0.0000 Office Killer (1997)  
+0.0000 War at Home, The (1996)  
+0.0000 Nosferatu a Venezia (1986)  
+0.0000 Killer: A Journal of Murder (1995)  
+0.0000 Sliding Doors (1998)

Bottom movies

-0.0000 Further Gesture, A (1996)  
-0.0000 Land and Freedom (Tierra y libertad) (1995)  
-0.0000 I, Worst of All (Yo, la peor de todas) (1990)  
-0.0000 Silence of the Palace, The (Saint el Qusur) (1994)  
-0.0000 Modern Affair, A (1995)  
-0.0000 Salut cousin! (1996)  
-0.0000 Wooden Man's Bride, The (Wu Kui) (1994)  
-0.0000 To Cross the Rubicon (1991)

\*\*\* User 2 \*\*\*

Top movies

+0.0000 Glass Shield, The (1994)  
+0.0000 Line King: Al Hirschfeld, The (1996)  
+0.0000 Sunchaser, The (1996)  
+0.0000 Land and Freedom (Tierra y libertad) (1995)  
+0.0000 Brothers in Trouble (1995)  
+0.0000 Golden Earrings (1947)  
+0.0000 Angela (1995)  
+0.0000 I, Worst of All (Yo, la peor de todas) (1990)

Bottom movies

-0.0000 Someone Else's America (1995)  
-0.0000 My Favorite Season (1993)  
-0.0000 Designated Mourner, The (1997)

-0.0000 Nothing Personal (1995)  
-0.0000 Last Klezmer: Leopold Kozlowski, His Life and Music, The  
(1995)  
-0.0000 Spirits of the Dead (Tre passi nel delirio) (1968)  
-0.0000 Death in Brunswick (1991)  
-0.0000 Foreign Student (1994)

\*\*\* User 3 \*\*\*

Top movies

+0.0005 Angel on My Shoulder (1946)  
+0.0004 Girls Town (1996)  
+0.0004 Mat' i syn (1997)  
+0.0003 Daens (1992)  
+0.0003 Eighth Day, The (1996)  
+0.0003 Lady of Burlesque (1943)  
+0.0003 Truth or Consequences, N.M. (1997)  
+0.0003 Silence of the Palace, The (Saint el Qusur) (1994)

Bottom movies

-0.0005 Leopard Son, The (1996)  
-0.0004 Salut cousin! (1996)  
-0.0003 Further Gesture, A (1996)  
-0.0003 Promise, The (Versprechen, Das) (1994)  
-0.0003 Men With Guns (1997)  
-0.0003 Woman in Question, The (1950)  
-0.0002 Foreign Student (1994)  
-0.0002 Invitation, The (Zaproszenie) (1986)

\*\*\* User 4 \*\*\*

Top movies

+0.0000 Target (1995)  
+0.0000 Nothing Personal (1995)  
+0.0000 Death in Brunswick (1991)  
+0.0000 Dadetown (1995)  
+0.0000 Brothers in Trouble (1995)  
+0.0000 Mat' i syn (1997)  
+0.0000 Land and Freedom (Tierra y libertad) (1995)  
+0.0000 Silence of the Palace, The (Saint el Qusur) (1994)

Bottom movies

-0.0000 Other Voices, Other Rooms (1997)  
-0.0000 Every Other Weekend (1990)  
-0.0000 Spanish Prisoner, The (1997)  
-0.0000 Wend Kuuni (God's Gift) (1982)  
-0.0000 Inkwell, The (1994)  
-0.0000 August (1996)  
-0.0000 Sleepover (1995)  
-0.0000 You So Crazy (1994)

## More Bonus Material: Interpretation of Weight Vectors as Features

- So far we have described both  $u_i$  and  $v_j$  as *weight vectors* (since we don't have any features of movies and users). But, it is possible to interpret one or both of these vectors as **learned features**.
- For example, the first learned feature may discover a preference for comedy vs. drama. In this case:
  - The user feature value  $u_{i1}$  should be high if the user likes comedies and low if the user likes dramas better.
  - The movie feature value  $v_{j1}$  should be high if the movie is a comedy and low if it is a drama.
- Similarly, feature 2 might describe whether a movie is geared toward kids or adults
- In practice, the feature interpretations often find recognizable patterns but are not quite so clean to describe as the two examples above.

Run the code below to examine the movies with the highest and lowest feature values for some of the features in your learned model.

```
k = 5
```

```
features_to_examine = np.arange(0,10)
```

```
for feature in features_to_examine:
```

```
    feature_vals = V[:,feature]
```

```
    print ('*** Feature %d ***' % (feature))
    print ('  Movies with highest feature value')
    for i in get_highest(feature_vals)[0:k]:
        print ('    %+.4f %s' % (feature_vals[i], titles[i]))
    print ('')
```

```
    print ('  Movies with lowest feature value')
    for i in get_lowest(feature_vals)[0:k]:
        print ('    %+.4f %s' % (feature_vals[i], titles[i]))
    print ('')
```

```
*** Feature 0 ***
```

```
Movies with highest feature value
+0.0270 Love and Death on Long Island (1997)
+0.0201 Crude Oasis, The (1995)
+0.0179 Woman in Question, The (1950)
+0.0171 Leopard Son, The (1996)
+0.0167 Girls Town (1996)
```

```
Movies with lowest feature value
-0.0218 Symphonie pastorale, La (1946)
```

- 0.0173 Silence of the Palace, The (Saint el Qusur) (1994)
- 0.0169 To Cross the Rubicon (1991)
- 0.0153 Further Gesture, A (1996)
- 0.0144 Jupiter's Wife (1994)

\*\*\* Feature 1 \*\*\*

Movies with highest feature value

- +0.0179 Jupiter's Wife (1994)
- +0.0166 Angela (1995)
- +0.0139 Tigrero: A Film That Was Never Made (1994)
- +0.0128 To Cross the Rubicon (1991)
- +0.0122 Death in the Garden (Mort en ce jardin, La) (1956)

Movies with lowest feature value

- 0.0249 Butcher Boy, The (1998)
- 0.0231 August (1996)
- 0.0213 Woman in Question, The (1950)
- 0.0201 Nemesis 2: Nebula (1995)
- 0.0155 Sleepover (1995)

\*\*\* Feature 2 \*\*\*

Movies with highest feature value

- +0.0216 Vermont Is For Lovers (1992)
- +0.0195 I, Worst of All (Yo, la peor de todas) (1990)
- +0.0166 Brothers in Trouble (1995)
- +0.0156 Sliding Doors (1998)
- +0.0154 Leopard Son, The (1996)

Movies with lowest feature value

- 0.0234 Hostile Intentions (1994)
- 0.0225 Star Kid (1997)
- 0.0218 Walk in the Sun, A (1945)
- 0.0182 Truth or Consequences, N.M. (1997)
- 0.0182 Target (1995)

\*\*\* Feature 3 \*\*\*

Movies with highest feature value

- +0.0288 Tigrero: A Film That Was Never Made (1994)
- +0.0172 Crude Oasis, The (1995)
- +0.0166 Dadetown (1995)
- +0.0159 Promise, The (Versprechen, Das) (1994)
- +0.0157 Someone Else's America (1995)

Movies with lowest feature value

- 0.0258 Wings of Courage (1995)
- 0.0219 Sliding Doors (1998)
- 0.0166 Brothers in Trouble (1995)
- 0.0164 Butcher Boy, The (1998)
- 0.0158 Line King: Al Hirschfeld, The (1996)

\*\*\* Feature 4 \*\*\*

Movies with highest feature value

+0.0250 Promise, The (Versprechen, Das) (1994)  
+0.0224 Woman in Question, The (1950)  
+0.0194 Wooden Man's Bride, The (Wu Kui) (1994)  
+0.0173 August (1996)  
+0.0161 Nightwatch (1997)

Movies with lowest feature value

-0.0277 Wend Kuuni (God's Gift) (1982)  
-0.0258 Last Klezmer: Leopold Kozłowski, His Life and Music, The  
(1995)  
-0.0233 Coldblooded (1995)  
-0.0195 Leopard Son, The (1996)  
-0.0134 Every Other Weekend (1990)

\*\*\* Feature 5 \*\*\*

Movies with highest feature value

+0.0203 Death in Brunswick (1991)  
+0.0189 Dadetown (1995)  
+0.0188 Venice/Venice (1992)  
+0.0151 Target (1995)  
+0.0143 Object of My Affection, The (1998)

Movies with lowest feature value

-0.0233 Big One, The (1997)  
-0.0220 Glass Shield, The (1994)  
-0.0220 Crude Oasis, The (1995)  
-0.0135 Wend Kuuni (God's Gift) (1982)  
-0.0124 Getting Away With Murder (1996)

\*\*\* Feature 6 \*\*\*

Movies with highest feature value

+0.0198 Last Klezmer: Leopold Kozłowski, His Life and Music, The  
(1995)  
+0.0166 Power 98 (1995)  
+0.0148 Every Other Weekend (1990)  
+0.0145 Daens (1992)  
+0.0144 Woman in Question, The (1950)

Movies with lowest feature value

-0.0262 Nothing Personal (1995)  
-0.0233 To Cross the Rubicon (1991)  
-0.0185 Shadows (Cienie) (1988)  
-0.0172 Salut cousin! (1996)  
-0.0156 Land and Freedom (Tierra y libertad) (1995)

\*\*\* Feature 7 \*\*\*

Movies with highest feature value

- +0.0259 King of New York (1990)
- +0.0216 8 Heads in a Duffel Bag (1997)
- +0.0168 Mat' i syn (1997)
- +0.0164 Power 98 (1995)
- +0.0157 Pushing Hands (1992)

Movies with lowest feature value

- 0.0219 Sliding Doors (1998)
- 0.0215 Other Voices, Other Rooms (1997)
- 0.0147 Angel on My Shoulder (1946)
- 0.0138 Wooden Man's Bride, The (Wu Kui) (1994)
- 0.0123 Butcher Boy, The (1998)

\*\*\* Feature 8 \*\*\*

Movies with highest feature value

- +0.0223 Lashou shentan (1992)
- +0.0213 Power 98 (1995)
- +0.0193 Shadows (Cienie) (1988)
- +0.0153 I Don't Want to Talk About It (De eso no se habla) (1993)
- +0.0141 Coldblooded (1995)

Movies with lowest feature value

- 0.0279 Modern Affair, A (1995)
- 0.0274 Walk in the Sun, A (1945)
- 0.0233 Dadetown (1995)
- 0.0206 Angel on My Shoulder (1946)
- 0.0189 Brothers in Trouble (1995)

\*\*\* Feature 9 \*\*\*

Movies with highest feature value

- +0.0227 Substance of Fire, The (1996)
- +0.0224 Getting Away With Murder (1996)
- +0.0184 Every Other Weekend (1990)
- +0.0155 Good Morning (1971)
- +0.0130 Nico Icon (1995)

Movies with lowest feature value

- 0.0213 Object of My Affection, The (1998)
- 0.0188 Brothers in Trouble (1995)
- 0.0184 Love and Death on Long Island (1997)
- 0.0175 Lady of Burlesque (1943)
- 0.0174 Lamerica (1994)